

Sapling E2E fixture: Lecture 3 notes on gradient descent.
Gradient descent iteratively updates parameters against the gradient of the loss with a fixed learning rate until the updates converge to a local minimum of the loss surface.